

Report Rubric Validation (plan-aligned)

The summer-project plan (目標二) requires a rubric checklist as the acceptance criterion, and the plan's §5.1 defines report validation as: covers the competition purpose + the five process components, is clear, reproducible-without-code, number-traceable, and honest.

The kaggle-report skill's own rubric (R1–R19) targets that skill's **7-section** format (e.g. R1 "7 節齊備", R17 "must not quote facts.json"), which is a *different* document structure from the 顏佐榕 **5-section** ML-spec framework used here. So these reports are validated against a **plan-aligned rubric** derived directly from Goal 2 and the plan's §5.1, not the skill's format-specific rubric.

Checklist (8 items)

#	Check	Standard
P1	Purpose — what	Overview + Purpose of Data state the problem being solved
P2	Purpose — why (重要性)	An explicit "Why it matters" states why the problem is important (non-canned)
P3	Five components complete	Data / Models / Training / Inference / Evaluation all present, each with substance or explicit "not recorded"
P4	Reproducible decisions	Splitting strategy + Feature Set + Reproducibility Standards documented so a data scientist could rebuild without code
P5	Numbers traceable	Every figure grounded in the competition's structured records (grounding line present; no invented numbers)
P6	Decision trajectory	Champion node identified and the score progression (stage/tier table) shown — the "breakthrough" trail
P7	Honesty	CV-only / LB-closed stated where true; unavailable items marked "not recorded"; source conflicts surfaced
P8	Performance benchmarking	Cross-agent comparison (vs NVIDIA reproduce-agent) present, supporting Goal 3

Result — all 25 validated reports pass 8/8

(15 original reports validated 2026-07-22; the 7 benchmark-baseline additions — afsis, cat-in-the-dat, conway, tps-aug-2022, tps-jan-2022, s5e1, tps-sep-2022 — drafted and number-audited 2026-08-17, every figure traced to the competition's workspace artifacts or the frozen three-way score table; the 3 special deep-learning-fallback lanes — ventilator, us-patent, siim-isic — added below, every figure traced to that competition's facts.json and the offline MLE-bench records in benchmark_results/run2/.)

Two of the rows below, **s3e20** and **s4e1**, sit outside the frozen baseline and their reports have been moved to ../archive_non_benchmark/; the md/ folder itself holds **23** reports — the 20 baseline competitions plus the 3 special lanes.

Comp	P1	P2	P3	P4	P5	P6	P7	P8
s3e1	✓	✓	✓	✓	✓	✓	✓	✓
s3e3	✓	✓	✓	✓	✓	✓	✓	✓
s3e5	✓	✓	✓	✓	✓	✓	✓	✓
s3e7	✓	✓	✓	✓	✓	✓	✓	✓
s3e9	✓	✓	✓	✓	✓	✓	✓	✓
s3e11	✓	✓	✓	✓	✓	✓	✓	✓
s3e14	✓	✓	✓	✓	✓	✓	✓	✓
s3e16	✓	✓	✓	✓	✓	✓	✓	✓
s3e19	✓	✓	✓	✓	✓	✓	✓	✓
s3e20	✓	✓	✓	✓	✓	✓	✓	✓
s4e1	✓	✓	✓	✓	✓	✓	✓	✓
s4e11	✓	✓	✓	✓	✓	✓	✓	✓
s5e10	✓	✓	✓	✓	✓	✓	✓	✓
s6e1	✓	✓	✓	✓	✓	✓	✓	✓
s6e2	✓	✓	✓	✓	✓	✓	✓	✓
afsis-soil-properties	✓	✓	✓	✓	✓	✓	✓	✓
cat-in-the-dat	✓	✓	✓	✓	✓	✓	✓	✓
conway-s-reverse-game-of-life	✓	✓	✓	✓	✓	✓	✓	✓
tps-aug-2022	✓	✓	✓	✓	✓	✓	✓	✓
tps-jan-2022	✓	✓	✓	✓	✓	✓	✓	✓
s5e1	✓	✓	✓	✓	✓	✓	✓	✓
tps-sep-2022	✓	✓	✓	✓	✓	✓	✓	✓
ventilator-pressure-prediction	✓	✓	✓	✓	✓	✓	✓	✓
us-patent-phrase-to-phrase-matching	✓	✓	✓	✓	✓	✓	✓	✓
siim-isic-melanoma-classification	✓	✓	✓	✓	✓	✓	✓	✓

For the three special lanes, P8 is satisfied by an offline MLE-bench three-way grade table (our agent vs NVIDIA vs AIDE) rather than a leaderboard comparison, and P7 additionally carries the "no live public/private leaderboard exists for this lane" statement.

Not covered by these reports (broader plan items, out of report scope): the plan's §5.2 code-consistency check (LLM compares report vs code), §5.3 reproducibility test (a third party rebuilds from report only), Goal 5 Docker containerization.