

model-level operators
seeded as starting nodes
(objective, blend member)

budget phase machine:
exploit → explore-burst → stop
(hard cap 60 evaluated nodes;
patience 20 evals without improvement)

on stop:
best node →
submission

select parent lineage
(backtrack when a lineage plateaus:
3 children without improvement)

priors from
experience library

propose mutation:
seed / hyperparameters /
blend weights / boundary push

duplicate gate (config hash; a repeated
proposal burns budget instead of looping)

data-level operators
each gets its own arm;
arms race, the winner
takes the full budget

evaluate in child process
(CV score; hard timeout; blend cost guard)

add node, update global best
(a solo breakthrough reopens
plateaued blend lineages)

loop