

AI Agents for Data Visualization and Kaggle Competitions

An end-to-end LLM agent (my-agent) benchmarked against two frozen yardsticks on 20 Kaggle competitions — with the apparatus needed to make the comparison honest

Author: Wei-Hao Huang (黄威昊) · **School:** National Chengchi University, Department of Economics
Advisor: Dr. Tso-Jung Yen (颜佐榕), Institute of Statistical Science, Academia Sinica · Summer Internship 2026 · Report v4 (2026-08-17)

1 · Motivation

Why these tasks are hard


- Kaggle competitions** demand the full expert loop — validation design, feature engineering, model selection, ensembling — where the difference between top entries sits at the 4th–5th decimal of the metric.
- Local cross-validation is an unreliable compass: **it contradicts the real leaderboard in 9 of 13** comparable competitions we measured.
- Data visualization** (EDA, diagnostics, reporting) is essential but repetitive: every dataset needs distribution, drift, and correlation views before modeling can be trusted.

Why an AI agent


- Tireless, systematic search: 22–80 candidate configurations per competition, every experiment logged and reproducible.
- Runs the same disciplined protocol on 20 competitions serially — a human cannot hold this consistency.
- Generates per-dataset EDA figures and per-competition reports automatically, so every claim is grounded in an artifact.

2 · Workflow

An LLM reasons at every stage; Auto-ML libraries are invoked through the shell; a cross-competition **experience library** supplies evidence-backed priors; a **tree search** over candidate configurations is the optimisation loop.



Six-stage pipeline: problem dossier → EDA → features → modeling → evaluation → submission. The v5 dossier layer classifies the task and injects external data *upstream*, before any modeling.



Tree-search loop: ensembles are first-class nodes with OOF caching and weight search; exploit until plateau, then one forced exploration burst.

4 · Advantages & Difficulties

Expected → experienced

- Expected:** the agent would quickly beat most humans. **Experienced:** development runs landed above average (median percentile 72.3), with wins concentrated where no mature public solutions exist; the frozen-architecture benchmark standing is TBD until scoring.
- Expected:** one run and one score decide a winner. **Experienced:** score noise + run-to-run variance leave **27% of duels undecidable**; a paired error model became mandatory.
- Expected:** the agent follows the rules by default. **Experienced:** a transcript audit caught one lane reading its own competition's recorded findings — and the clean re-run scored *worse*, proving the contamination advantage was real.
- Fairness is expensive: isolated data roots, credential-free sandboxes, and per-lane transcript audits cost real engineering time — but without them the numbers are not comparable.


3 · Results — Kaggle benchmark

All submissions scored on the real leaderboard (late submission); verdicts adjudicated by a paired test built from Kaggle's own public/private split.

Benchmark standing: TBD until scoring

All 20 competitions run under the frozen Stage-0.5 architecture in credential-free sandboxes — submissions exist, leaderboard scoring is a separate credentialed step. Every verdict is then adjudicated by a paired test built from Kaggle's own public/private split.

Scored today: the injection-layer ablation (s3e19, private SMAPE ↓)



The dossier layer's external-data injection turns the architecture's canonical failure into a numerical first among the three agents — and all 6 of 6 controlled external-data arms beat their matched baselines. Bar length ∝ margin over 52 (lower SMAPE is better).


Mid-complexity three-way (offline-graded, private split)

Competition	Metric	my-agent	NVIDIA	AIDE
us-patent (NLP)	Pearson r ↑	0.86658	0.86160	0.65821
ventilator (time series)	MAE ↓	0.16525	0.40093	0.34591
SIIM-ISIC (imaging)	AUC ↑	0.93529	0.89822	0.88579

Final-architecture re-run (in progress, 2026-08-17)

All 20 competitions re-run under one frozen architecture, credential-free sandboxes, per-lane transcript audit. **16/20 lanes run** (one pending audit review), 1 running, 3 queued — leaderboard scores **TBD** until the operator scoring step.

3b · Results — visualization output



Agent-generated EDA: target distribution and temporal-pattern diagnostics, produced automatically per dataset before modeling (example: energy-prosumer competition). The same pipeline emits per-competition ML-spec reports (22 published) and the flowcharts above.

5 · Discussion — advice

- State the error model with every claim.** A ranking without an error scale is noise: 27% of our duels are undecidable, and the winner flips when the evaluation set changes.
- Never trust local CV for between-agent claims.** Score on the real leaderboard; CV describes a pipeline's internal view only.
- Freeze your yardsticks.** Improving a baseline mid-study silently turns the benchmark into an optimisation of the baseline.
- Isolate mechanically, then audit the transcript.** Sandboxes stop leaks you anticipated; reading what the agent actually did catches the ones you didn't — in both directions (we hit one real contamination and one auditor false positive).
- Put knowledge injection upstream.** Injected at the blend stage it was worth $\sim 10^{-5}$; moved into problem identification it lifted the canonical failure case from last to first.

Where knowledge enters

- Downstream (blend stage): measured gain $\sim 10^{-5}$ — abandoned.
- Upstream (problem dossier): fires on exactly the 2/18 competitions that need external data; all **6 of 6** controlled external-data arms beat their matched baselines.
- But no local signal can select the *form* of injection — the one surviving selector died on its first pre-registered prediction.

Companion documents: Report v4 (Overleaf), Engineering Log, Case Studies, 22 per-competition ML-spec reports.